

PolicyImpact

Evidence-grounded policy change-to-claim review

Jiuyi Zheng · Synthetic policies and claims · No PHI · No production decisions

The opportunity is governed interpretation—not autonomous adjudication

WHY NOW

Public healthcare guidance increasingly emphasizes transparent decision support, data provenance, and governed AI. The tractable product surface is the handoff from policy evidence to human-reviewed rules—not replacing domain experts.

Opportunity

- Evidence-linked change triage
- Typed rule proposals
- Expert review queues

Threat

- Unsupported rule fields
- Silent automation
- Scope drift into adjudication

Guardrail: generative output stays non-executable until Auditor review and explicit human approval.

Architecture keeps generative roles outside execution

Interpret

Versioned policy pair
Controlled parser
Analyst proposal

Govern

Independent Auditor
Exactly one revision
Explicit human gate

Execute

Approved typed rules
Deterministic engine
Exact evidence trace

01 SOURCE

02 GATE

03 REVIEW QUEUE

Two roles, one revision, one explicit decision boundary

Ambiguity remains visible and non-executable; abstention is an allowed outcome.

ROUND 1

Analyst proposes

4 changes
Typed candidate pairs
Exact sections + caveats

CHALLENGE

Auditor tests


Accept / revise / abstain
Modifier logic challenged
No hidden rationale

ROUND 2

Revision closes

3 accept · 1 abstain
Human approve / reject
No second AI loop

Controlled fixture replay validates the workflow plumbing



OFFLINE FIXTURE REPLAY

All reported values are deterministic checks against synthetic gold fixtures. They are not live-model accuracy, generalization, production savings, or clinical/payment performance.

16 / 16

claim-impact records match
gold

20 / 20

outputs retain exact
traceability

Final review: 3 accepted · 1 abstained · 0 unsupported assertions

Recommendation: run a result-blind, multi-policy pilot

The POC is ready to test as a governed review aid—not to deploy as an autonomous decision-maker.

BOUNDARIES

**No PHI · no clinical advice · no payment decision
· no production deployment · no savings claim**

- ✓ Freeze 10–20 approved policy pairs and gold criteria before model runs
- ✓ Use independent domain experts to author and adjudicate ground truth
- ✓ Run in a controlled environment; measure accuracy, abstention, traceability, and overrides
- ✓ Stop or roll back when evidence, safety, or approval gates fail